

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

dataset = pd.read_csv("diabetes.csv")
dataset
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

```
[768 rows x 9 columns]
```

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

```
dtypes: float64(2), int64(7)
```

```
memory usage: 54.1 KB
```

```
dataset.isnull().sum()
```

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0

```
dtype: int64
```


```
dataset.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458
std	3.369578	31.972618	19.355807	15.952218
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000
75%	6.000000	140.250000	80.000000	32.000000
max	17.000000	199.000000	122.000000	99.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000


```
plt.figure(figsize=(10,8))
sns.heatmap(dataset.corr(),annot = True, fmt=".3f",cmap="YlGnBu")
plt.title("Correlation Heatmap")
```

```
Text(0.5, 1.0, 'Correlation Heatmap')
```




```
plt.figure(figsize=(10,8))
kde = sns.kdeplot(dataset["Pregnancies"]
[dataset["Outcome"]==1],color="red",fill=True)
kde = sns.kdeplot(dataset["Pregnancies"]
[dataset["Outcome"]==0],color="blue",fill=True)
kde.set_xlabel("Pregnancies")
kde.set_ylabel("Density")
kde.legend(["Positive","Negatives"])

<matplotlib.legend.Legend at 0x162d705a5c0>
```




```
plt.figure(figsize=(10,8))
sns.violinplot(data=dataset,x="Outcome",y="Glucose",split=True,linewidth
th=2,inner="quart")

<Axes: xlabel='Outcome', ylabel='Glucose'>
```



```
plt.figure(figsize=(10,8))
kde = sns.kdeplot(dataset["Glucose"]
[dataset["Outcome"]==1],color="red",fill=True)
kde = sns.kdeplot(dataset["Glucose"]
[dataset["Outcome"]==0],color="blue",fill=True)
kde.set_xlabel("Glucose")
kde.set_ylabel("Density")
kde.legend(["Positive","Negatives"])
<matplotlib.legend.Legend at 0x162d7200310>
```



```
dataset["Glucose"]=dataset["Glucose"].replace(0,dataset["Glucose"].median())
dataset["BloodPressure"]=dataset["BloodPressure"].replace(0,dataset["BloodPressure"].median())
dataset["BMI"]=dataset["BMI"].replace(0,dataset["BMI"].mean())
dataset["SkinThickness"]=dataset["SkinThickness"].replace(0,dataset["SkinThickness"].mean())
dataset["Insulin"]=dataset["Insulin"].replace(0,dataset["Insulin"].mean())
```

dataset

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
BMI \ 0	6	148	72	35.000000	79.799479
33.6					
1	1	85	66	29.000000	79.799479
26.6					
2	8	183	64	20.536458	79.799479
23.3					

3	1	89	66	23.000000	94.000000
28.1					
4	0	137	40	35.000000	168.000000
43.1					
..
...					
763	10	101	76	48.000000	180.000000
32.9					
764	2	122	70	27.000000	79.799479
36.8					
765	5	121	72	23.000000	112.000000
26.2					
766	1	126	60	20.536458	79.799479
30.1					
767	1	93	70	31.000000	79.799479
30.4					

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

X=dataset.drop(["Outcome"],axis=1)

Y=dataset["Outcome"]

X

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
BMI \					
0	6	148	72	35.000000	79.799479
33.6					
1	1	85	66	29.000000	79.799479
26.6					
2	8	183	64	20.536458	79.799479
23.3					
3	1	89	66	23.000000	94.000000
28.1					
4	0	137	40	35.000000	168.000000
43.1					
..

```

...
763      10      101      76      48.000000      180.000000
32.9
764       2      122      70      27.000000      79.799479
36.8
765       5      121      72      23.000000      112.000000
26.2
766       1      126      60      20.536458      79.799479
30.1
767       1       93      70      31.000000      79.799479
30.4

```

```

      DiabetesPedigreeFunction  Age
0                0.627      50
1                0.351      31
2                0.672      32
3                0.167      21
4                2.288      33
..
763              0.171      63
764              0.340      27
765              0.245      30
766              0.349      47
767              0.315      23

```

[768 rows x 8 columns]

Y

```

0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0

```

Name: Outcome, Length: 768, dtype: int64

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.33,random_state=42)
```

x_train

```

      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin
BMI \

```


464	10	115	98	20.536458	79.799479
24.0					
223	7	142	60	33.000000	190.000000
28.8					
393	4	116	72	12.000000	87.000000
22.1					
766	1	126	60	20.536458	79.799479
30.1					
570	3	78	70	20.536458	79.799479
32.5					
..
...					
71	5	139	64	35.000000	140.000000
28.6					
106	1	96	122	20.536458	79.799479
22.4					
270	10	101	86	37.000000	79.799479
45.6					
435	0	141	72	20.536458	79.799479
42.4					
102	0	125	96	20.536458	79.799479
22.5					

	DiabetesPedigreeFunction	Age
464	1.022	34
223	0.687	61
393	0.463	37
766	0.349	47
570	0.270	39
..
71	0.411	26
106	0.207	27
270	1.136	38
435	0.205	29
102	0.262	21

[514 rows x 8 columns]

```

from sklearn.neighbors import KNeighborsClassifier


train_acc=[]
test_acc=[]
for n_neighbors in range(1,11):
    knn=KNeighborsClassifier(n_neighbors=n_neighbors)
    knn.fit(x_train,y_train)
    train_acc.append(knn.score(x_train,y_train))
    test_acc.append(knn.score(x_test,y_test))

plt.plot(range(1,11),train_acc,label="Train_acc")
plt.plot(range(1,11),test_acc,label="Test_acc")

```

```
plt.ylabel("Accuracy")
plt.xlabel("n_neighbors")
plt.legend()
```

<matplotlib.legend.Legend at 0x162e0dcbac0>



```
knn=KNeighborsClassifier(n_neighbors=9)
knn.fit(x_train,y_train)
print(knn.score(x_train,y_train),": Training accuracy")
print(knn.score(x_test,y_test),": Test accuracy")
```

```
0.7879377431906615 : Training accuracy
0.7244094488188977 : Test accuracy
```

```
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(random_state=0)
dt.fit(x_train,y_train)
print(dt.score(x_train,y_train),": Training accuracy")
print(dt.score(x_test,y_test),": Test accuracy")
```

```
1.0 : Training accuracy
0.6811023622047244 : Test accuracy
```

```

dt1=DecisionTreeClassifier(random_state=0,max_depth=3)
dt1.fit(x_train,y_train)
print(dt1.score(x_train,y_train),": Training accuracy")
print(dt1.score(x_test,y_test),": Test accuracy")

0.77431906614786 : Training accuracy
0.6929133858267716 : Test accuracy

from sklearn.neural_network import MLPClassifier
mlp=MLPClassifier(random_state=42)
mlp.fit(x_train,y_train)
print(mlp.score(x_train,y_train),": Training accuracy")
print(mlp.score(x_test,y_test),": Test accuracy")

0.7509727626459144 : Training accuracy
0.6811023622047244 : Test accuracy

from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train_sc=sc.fit_transform(x_train)
x_test_sc=sc.fit_transform(x_test)

mlp1=MLPClassifier(random_state=42)
mlp1.fit(x_train_sc,y_train)
print(mlp1.score(x_train_sc,y_train),": Training accuracy")
print(mlp1.score(x_test_sc,y_test),": Test accuracy")

0.8346303501945526 : Training accuracy
0.7362204724409449 : Test accuracy

C:\Users\sujan\anaconda3\lib\site-packages\sklearn\neural_network\
_multilayer_perceptron.py:684: ConvergenceWarning: Stochastic
Optimizer: Maximum iterations (200) reached and the optimization
hasn't converged yet.
  warnings.warn(

```